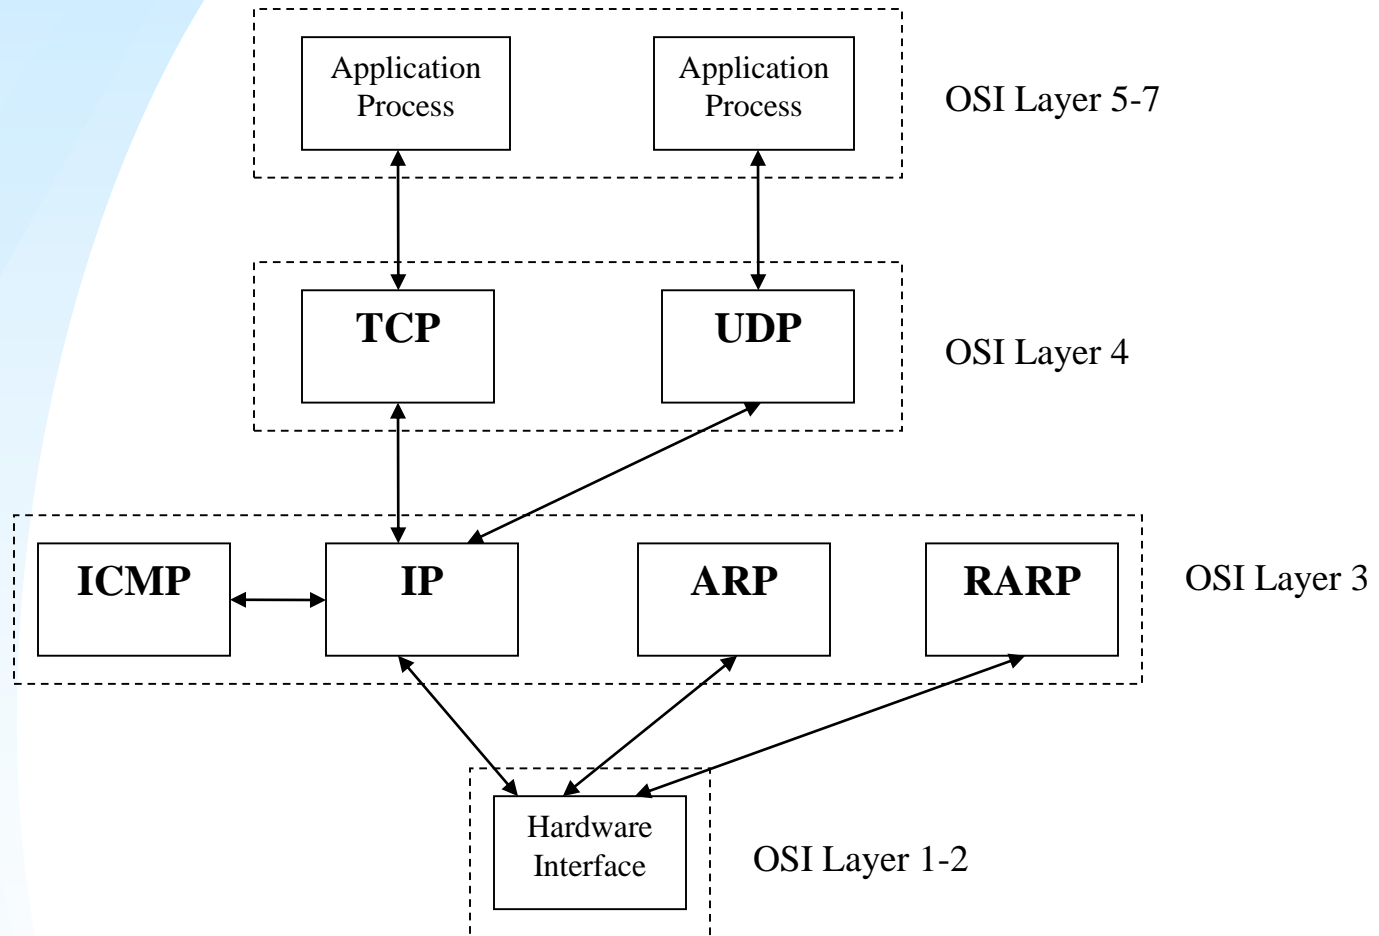


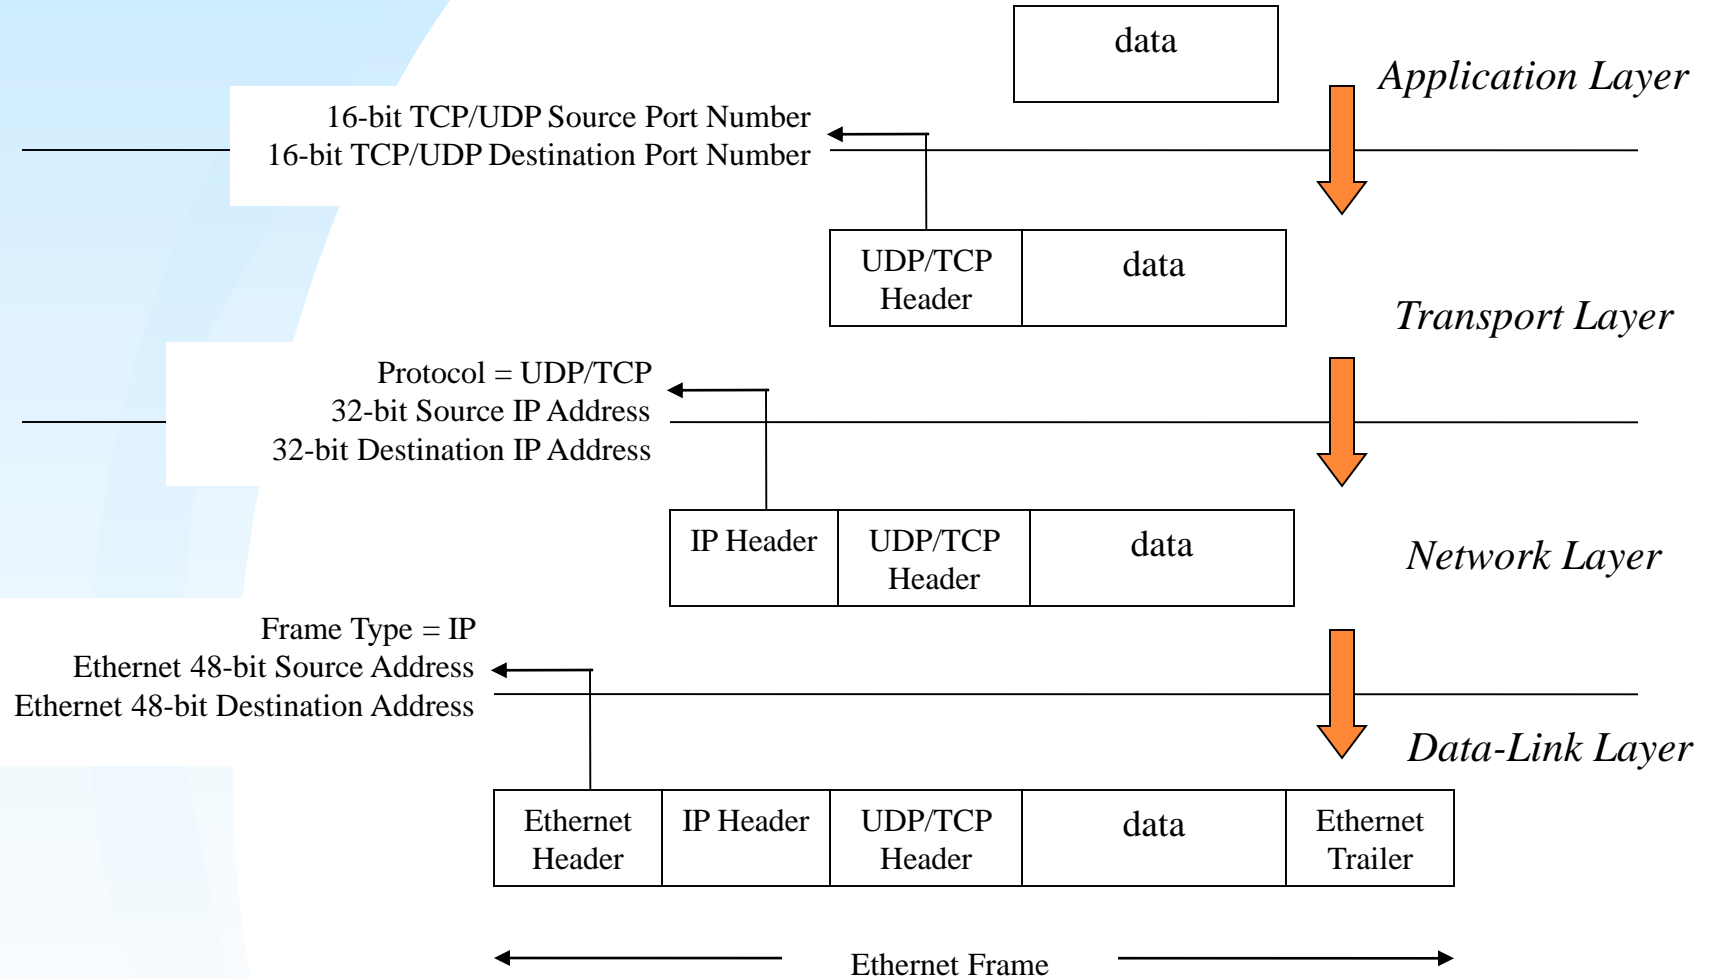
# ***TCP/IP Socket Programming in C***

**Dr. Sujoy Saha**  
**Assistant Professor**  
**NIT Durgapur**

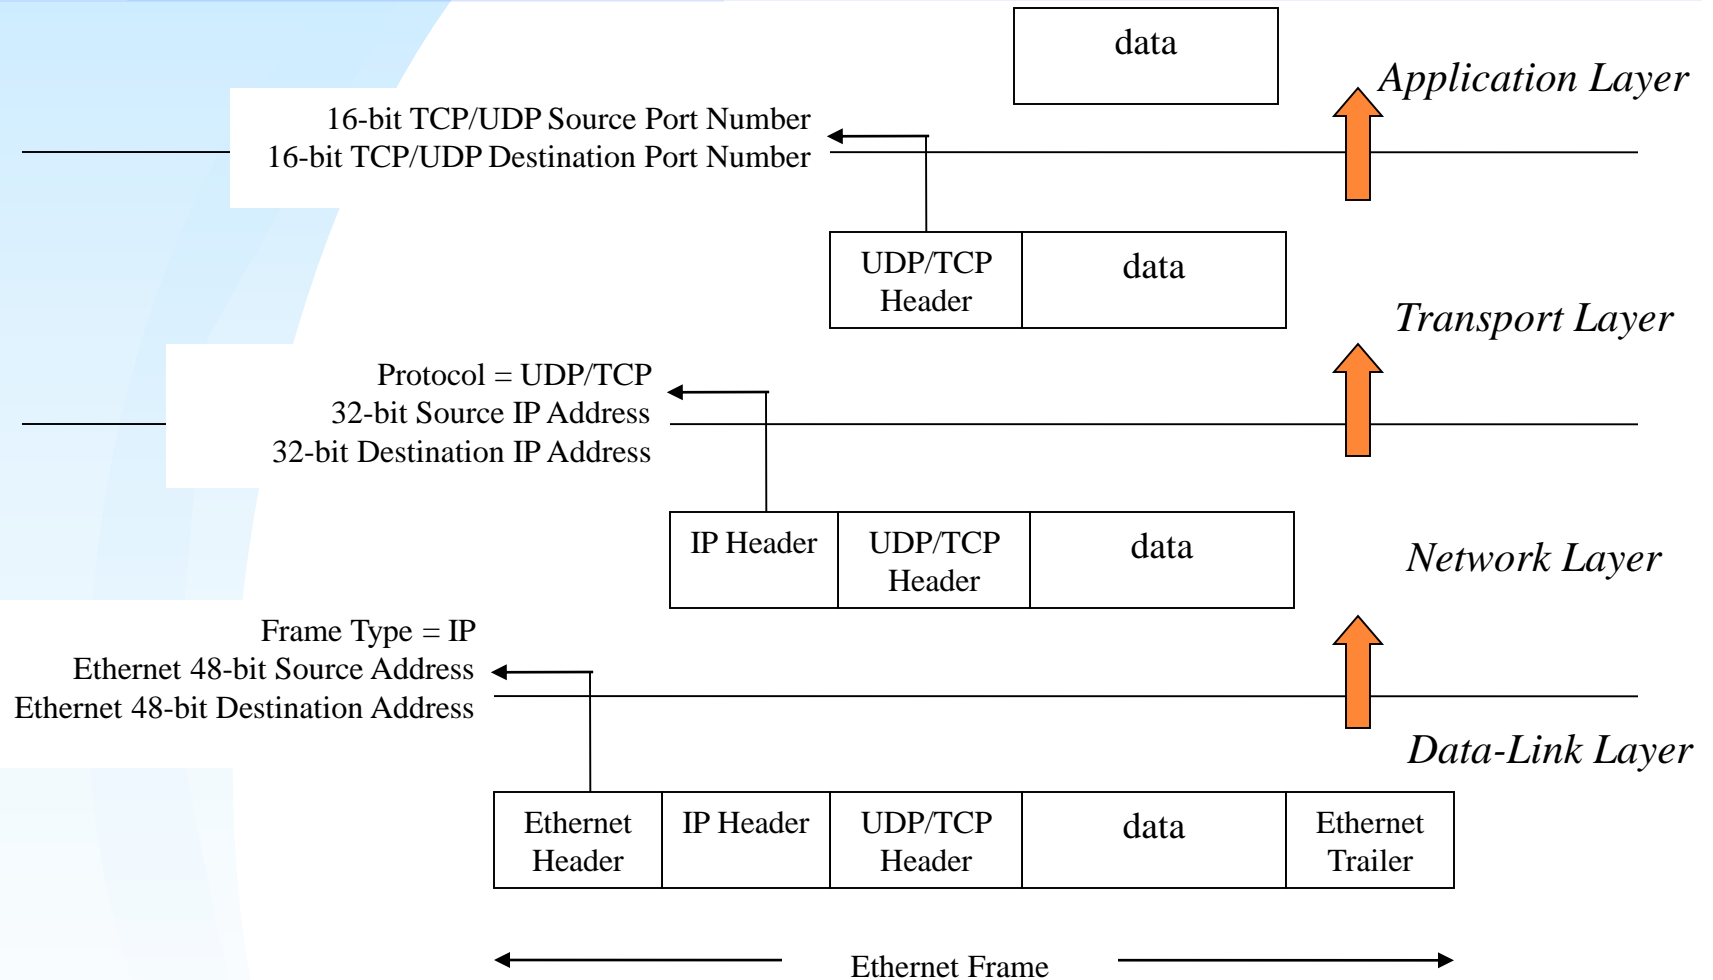
# TCP/IP Protocol Suite



# Sending Data



# Receiving Data




# Connecting Network with Operating System

---

# Socket

“A **network socket** is an endpoint of an inter-process communication across a computer network.”



# Connection

- In TCP/IP protocol suite
  - a connection defines the communication link between two processes
  - An *association* defined by the 5-tuple completely specifies two processes that make up a connection

*{protocol, local-addr, local-port, foreign-addr, foreign-port}*

- The protocol (TCP or UDP)
- The local host's IP address (32-bit)
- The local port number (16-bit)
- The foreign host's IP address (32-bit)
- The foreign port number (16-bit)

- Example

`{tcp, 192.168.2.2, 1500, 192.168.2.10, 21}`

# Sockets

- Sockets are *application program interface (API)* to the communication protocol.
- Availability of an API depends both on the operating system and the programming language.



# Socket Addresses

- For Internet family, following structures are defined in

`<netinet/in.h>`

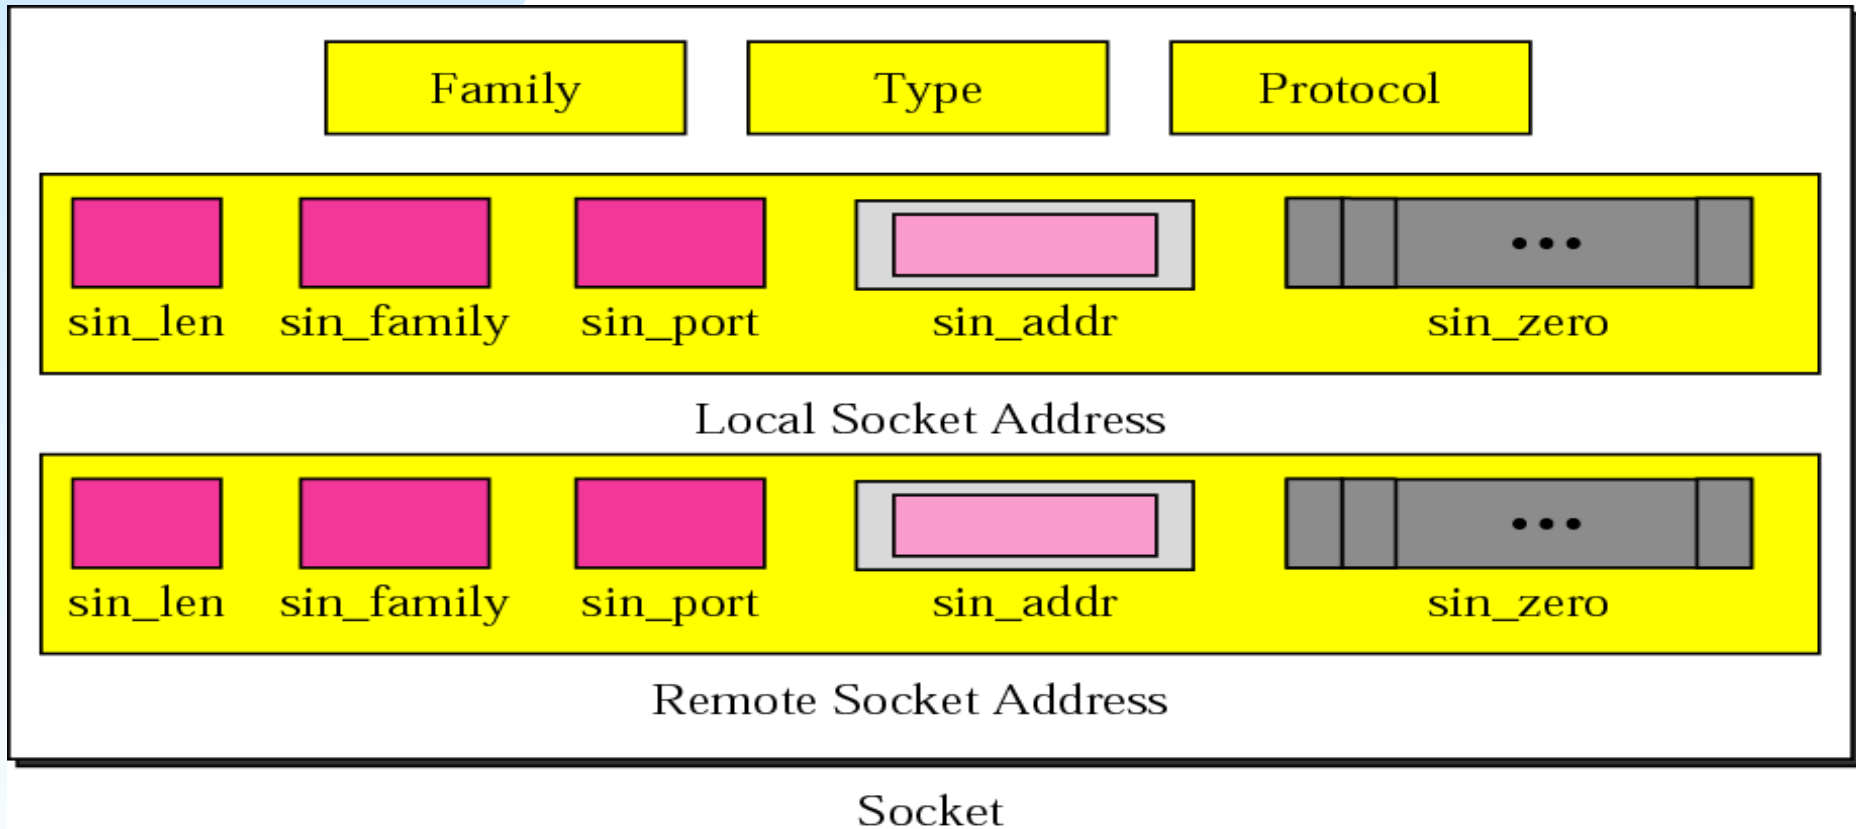
```
struct in_addr {  
    u_long      s_addr;      /*32-bit netid/hostid*/  
};
```

```
struct sockaddr_in {  
    short      sin_family; /*AF_INET*/  
    u_short    sin_port;   /*16-bit port number*/  
    struct in_addr sin_addr; /*32-bit netid/hostid*/  
    char       sin_zero[8]; /*unused*/  
};
```

- The header file `<sys/types.h>` provides data type definitions

Figure 16-4

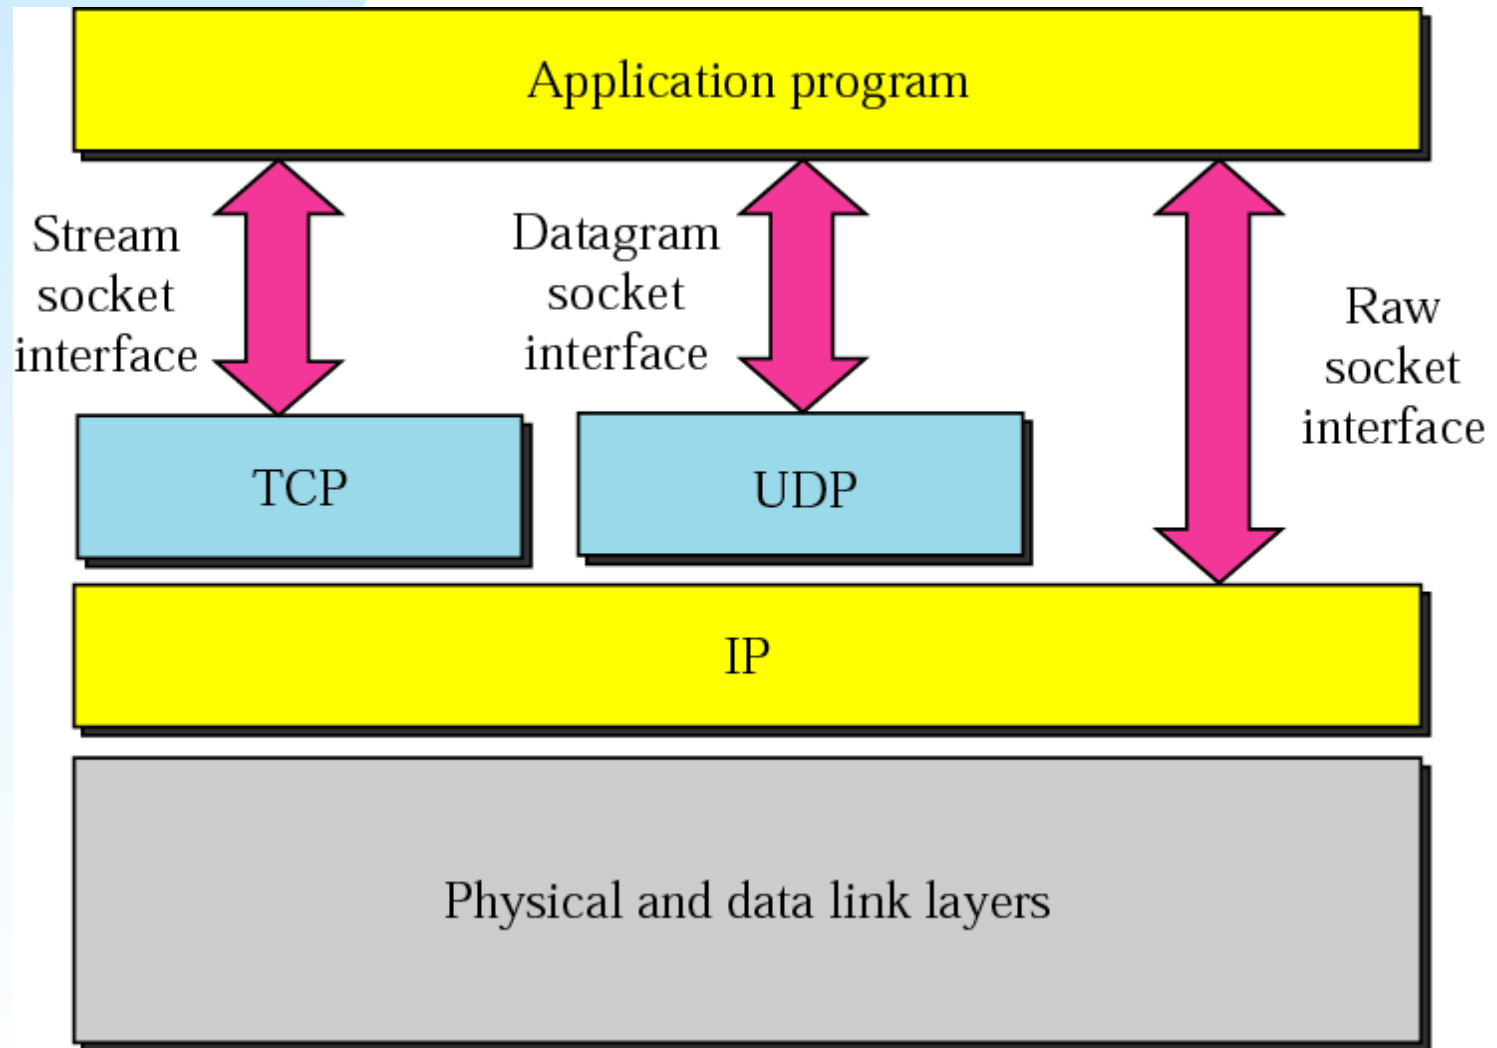
# Socket structure



Communication domain in which the socket should be created. Some of address families are `AF_INET` (IP), `AF_INET6` (IPv6), `AF_UNIX` (local channel, similar to pipes), `AF_ISO` (ISO protocols), and `AF_NS` (Xerox Network Systems protocols).

Figure 16-5

# Socket types





# **BYTE ORDERING**

Figure 16-6

## Big-endian byte order

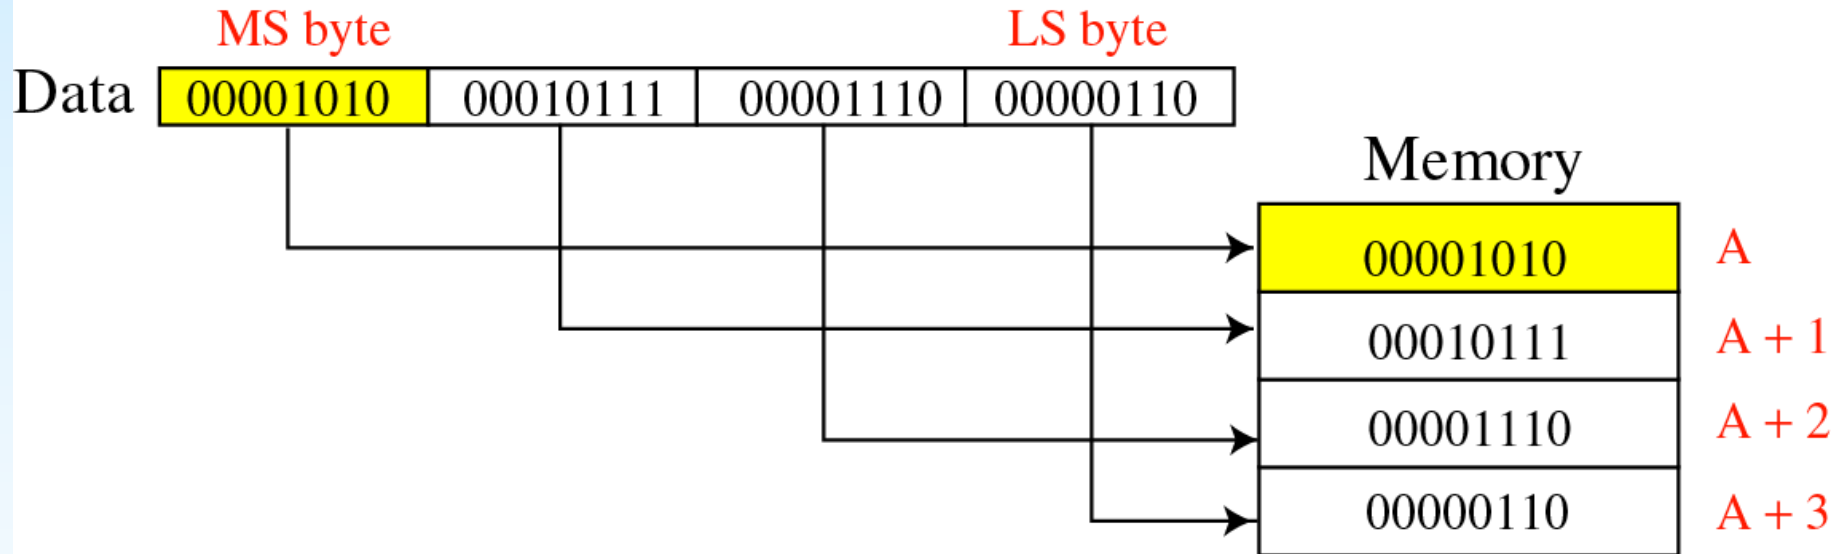
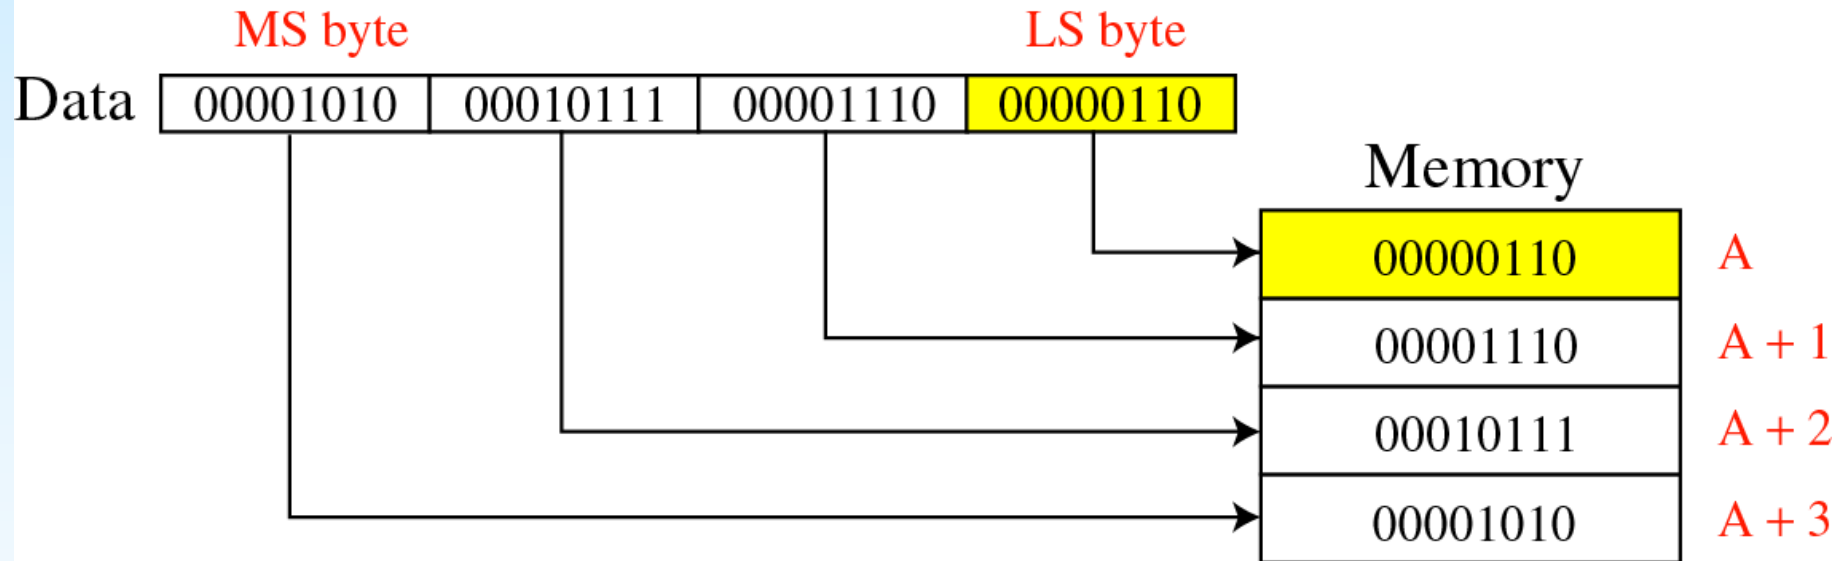


Figure 16-7

## Little-endian byte order



Note

*The byte order for the TCP/IP protocol suite is big endian.*

# Byte Ordering Functions


- Hosts can store multi-byte values differently (*host byte order*)
  - Little-endian byte order
  - Big-endian byte order
- Network protocols must specify a *network byte order*
  - Internet protocols use big-endian byte ordering
- Byte ordering routines

```
u_long htonl(u_long hostlong);  
u_short htons(u_short hostshort);  
u_long ntohl(u_long netlong);  
u_short ntohs(u_short netshort);
```



Figure 16-8

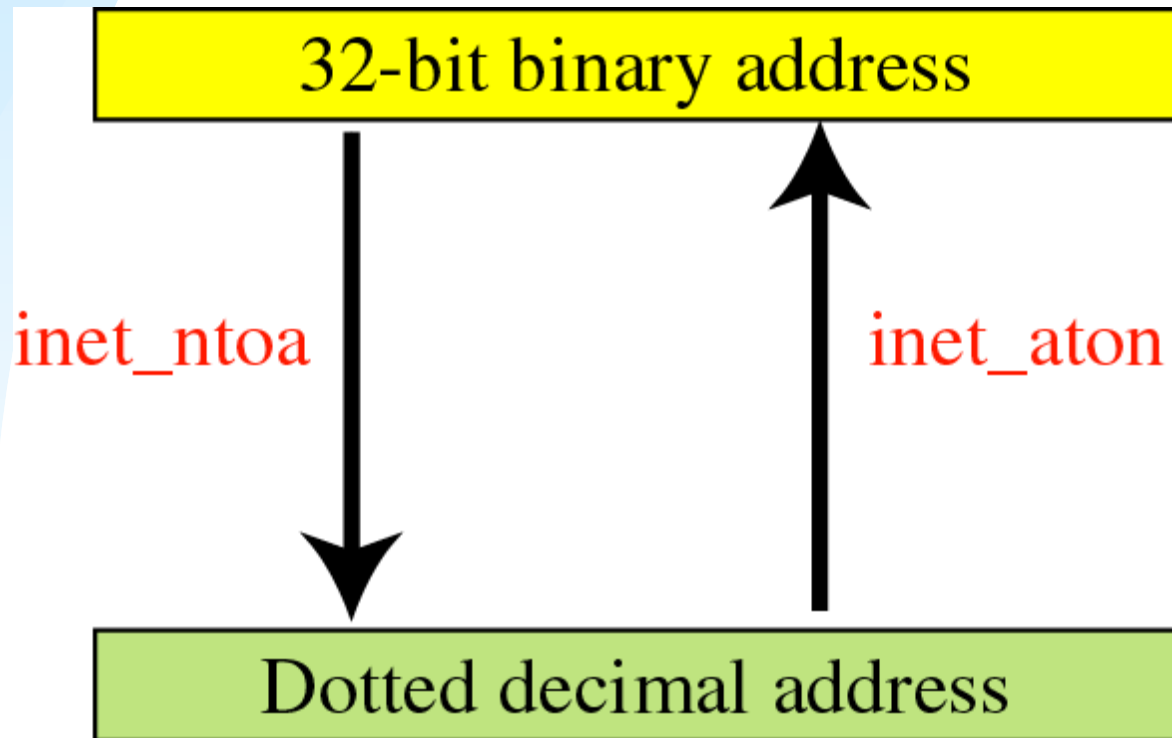
# Bite-order transformation



# Data types

<b>int8_t</b>	Signed 8-bit integer
<b>int16_t</b>	Signed 16-bit integer
<b>int32_t</b>	Signed 32-bit integer
<b>uint8_t</b>	Unsigned 8-bit integer
<b>uint16_t</b>	Unsigned 16-bit integer
<b>uint32_t</b>	Unsigned 32-bit integer

## Address transformation



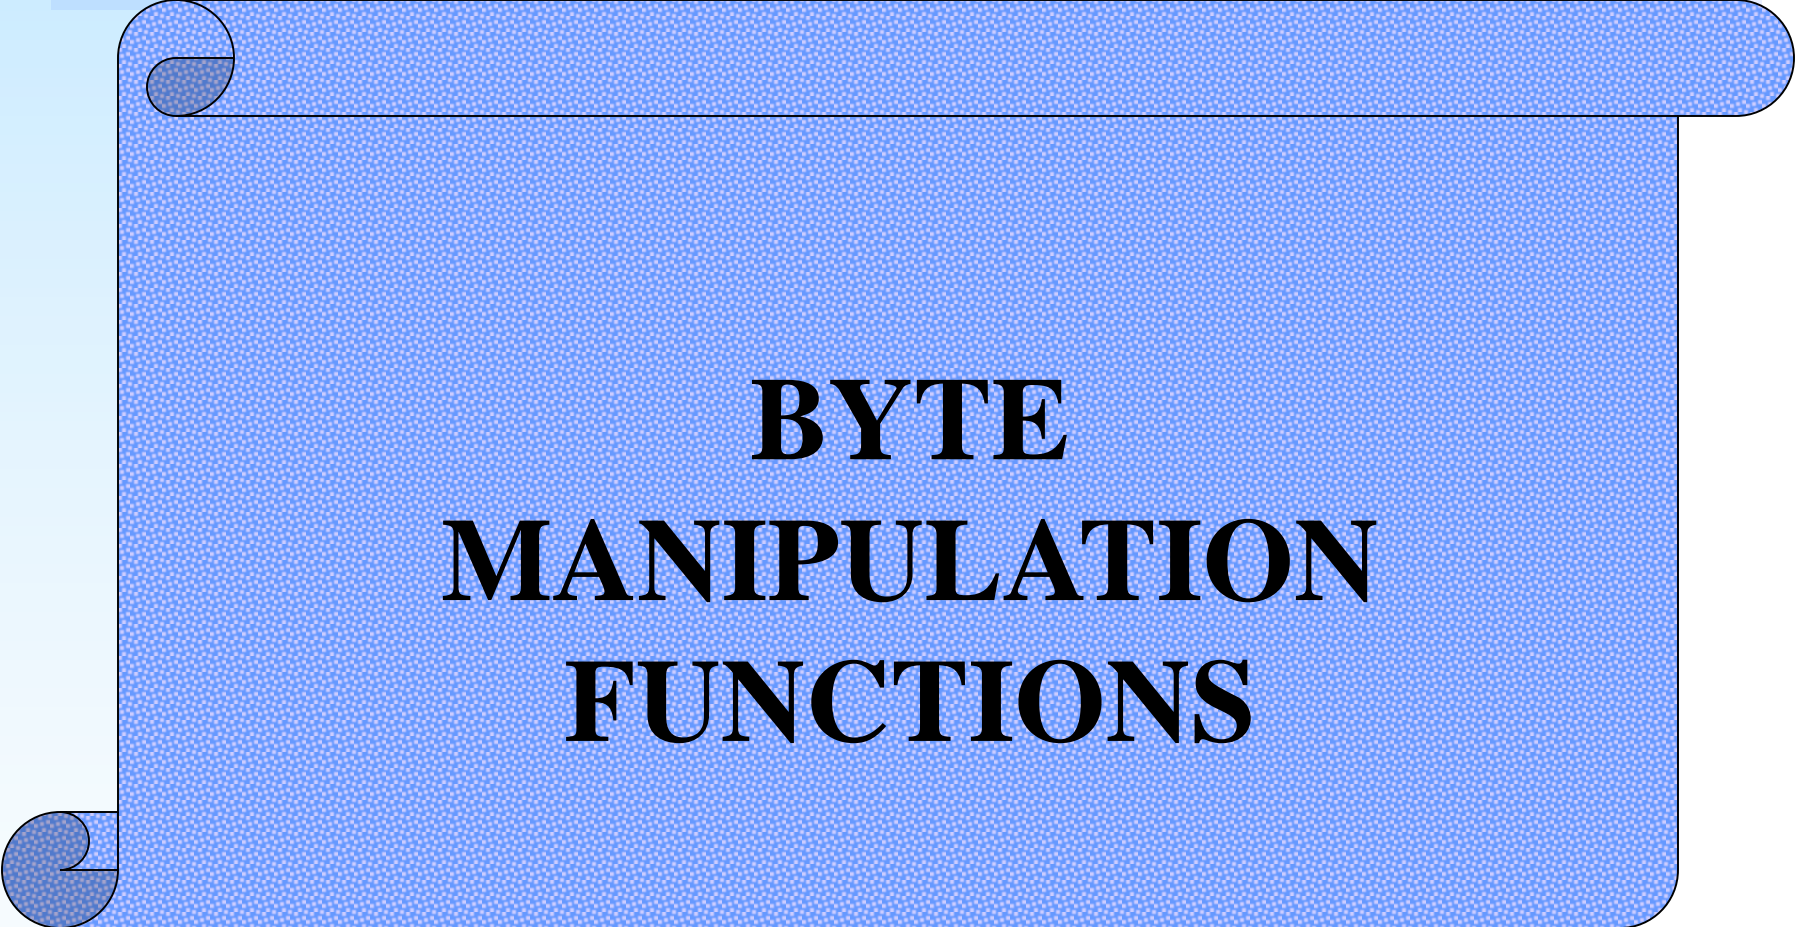
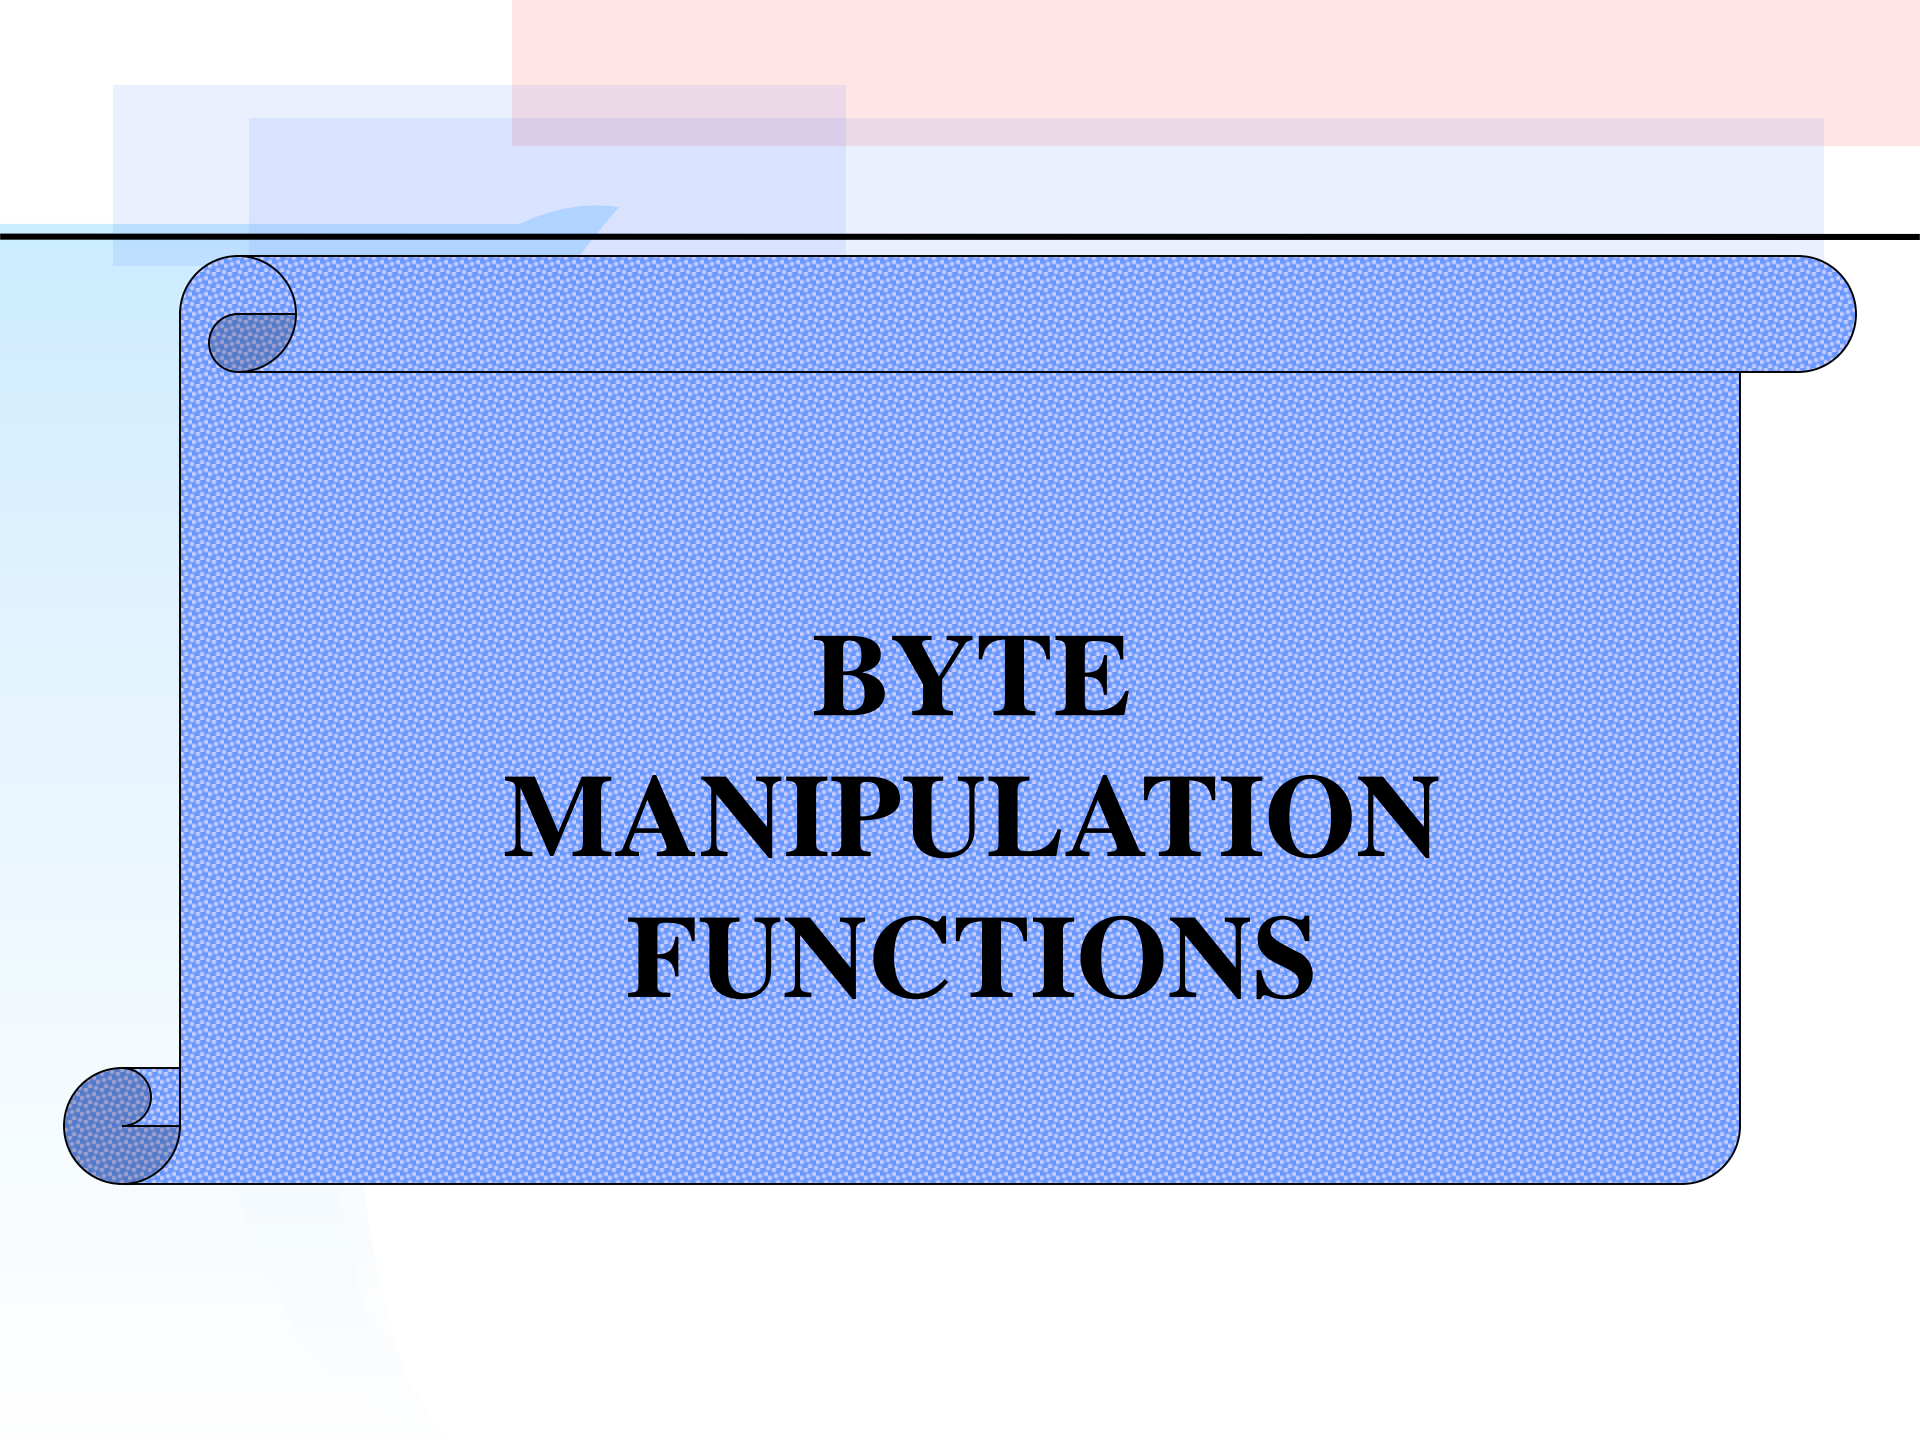
# Address Conversion Routines

```
unsigned long inet_addr(char *ptr);  
int inet_aton(char *ptr, struct in_addr *addrptr);
```

Converts a character string in dotted-decimal notation to a 32-bit Internet address in network byte order

```
char *inet_ntoa(struct in_addr inaddr);
```

Converts a 32-bit Internet address in network byte order to a character string in dotted decimal notation



# **BYTE MANIPULATION FUNCTIONS**

# Declarations for byte-manipulation functions

```
void      *memset ( void *dest , int  chr , size_t len ) ;
```

```
void      *memcpy ( void *dest , const void  *src , size_t len ) ;
```

```
int        memcmp ( const void *first , const void  *second , size_t len ) ;
```

```
char buffer[1400];
```

```
memset(buffer,0,sizeof(buffer));
```

```
bzero(buffer, sizeof(buf));
```

```
memcpy(buffer, msg.m_txt, sizeof(msg.m_txt));
```

```
struct MSG
{
    int m_txt[MAX_TEXT];
    int m_seq;
    int
    total_no_of_packets;
};
```



# **SOCKET SYSTEM CALLS**

# socket System Call

```
int socket(int family, int type, int protocol);
```

- The family is one of

AF_UNIX	Unix internal protocols
AF_INET	Internet Protocols
AF_NS	Xerox NS Protocols
AF_IMPLINK	IMP link Layer

- AF\_INET uses the TCP/IP protocol.
- AF\_UNIX creates filesystem objects and it only works between processes on the same host.
- AF\_NS is a set of protocols that were used by Xerox Systems for data communication. Its basic **working mechanism** is almost the **same as in the TCP/IP protocol** suit, but **XNS contains only two network layers**



# AF\_INET and AF\_UNIX

- AF\_INET uses the TCP/IP protocol.
- AF\_UNIX creates filesystem objects and it only works between processes on the same host.
- AF\_UNIX is much faster than AF\_INET.
- Xerox Network Systems (XNS) is a set of protocols that were used by Xerox Systems for data communication. Xerox used **XNS for file transfers, sharing network resources, packet transfers, sharing routing information and remote procedure calls**. Its basic **working mechanism** is almost the **same as in the TCP/IP protocol** suit, but **XNS contains only two network layers**. This differs from the seven-layer Open Systems Interconnection (OSI) model, although the functionality is basically the same.

# socket System Call

```
int socket(int family, int type, int protocol);
```

- The family is one of

AF_UNIX	Unix internal protocols
AF_INET	Internet Protocols
AF_NS	Xerox NS Protocols
AF_IMPLINK	IMP link Layer

- The socket type is one of the following

SOCK_STREAM	Stream Socket
SOCK_DGRAM	Datagram Socket
SOCK_RAW	Raw Sockets
SOCK_SEQPACKET	Sequenced Packet Socket
SOCK_RDM	Reliably Delivered Message Socket

- The protocol argument to the socket system call is typically set to 0
- The `socket` system call returns an integer value called socket descriptor
- This number is passed as a parameter to almost all of the other library calls.

Figure 16-15

# Declaration for *socket* function

```
int socket (int family , int type , int protocol) ;
```

↑                      ↑                      ↑  
 AF\_INET    SOCK\_DGRAM    0  
              SOCK\_STREAM

Returns a socket descriptor if successful; -1 if error.



- The `socket()` system call returns a file descriptor.
- In fact, a socket is similar to an opened file because it is possible to read and write data on it by means of the usual `read()` and `write()` system calls.
- in the context of the client-server model, the socket system calls needed to create and connect a pair of sockets and transmit data.

<i>type</i> \ <i>family</i>	AF_INET
SOCK_STREAM	<b>TCP</b> or SCTP
SOCK_DGRAM	<b>UDP</b>
SOCK_SEQPACKET	SCTP
SOCK_RAW	IPv4

<i>Protocol</i>
IPPROTO_TCP
IPPROTO_UDP
IPPROTO_SCTP



## Family

**AF\_UNIX**

**AF\_INET**

**AF\_ISO**

**AF\_NS**

**AF\_IPX**

**AF\_APPLETALK**

## Description

**UNIX internal (file system sockets)**

**ARPA Internet protocols (UNIX network sockets)**

**ISO standard protocols**

**Xerox Network Systems protocols**

**Novell IPX protocol**

**Appletalk DDS**

# bind

## bind system call:

- `bind()` associates the socket with its local address
- The `bind()` system call **binds an address**

```
#include <sys/socket.h>
int bind( int sockfd, const struct sockaddr * myaddr, socklen_t
addrlen );
```

*Returns: 0 if OK, -1 on error*

sockaddr is a generic structure, bind system call uses typepunning

### ■ e.g. `bind()`

```
myaddr.sin_family = AF_INET;
myaddr.sin_addr.s_addr = htonl( INADDR_ANY );
myaddr.sin_port = htons( 7 ); // 7 for echo service
```

```
if( bind( sockfd, (struct sockaddr *) &myaddr,
sizeof(myaddr) ) < 0 ) { ... }
```

# bind System Call

- Uses of `bind` system call
  - Servers register their well-known address with the system; both connection-oriented and connection-less servers need to do this
  - A client can register a specific address for itself
- The `bind` system call specifies the *local-addr* and *local-port* elements of the 5-tuple of an association

```
struct sockaddr_in {
    short    sin_family; /*AF_INET*/
    u_short  sin_port;   /*16-bit port number*/
    struct in_addr sin_addr; /*32-bit netid/hostid*/
    char     sin_zero[8]; /*unused*/
};
```

```
Struct sockaddr
{
    unsigned short sa_family;
    Char sa_data[14]
}
```

$\text{sin\_port (2 Byte)} + \text{Sin\_addr(4 Byte)} + \text{sin\_Zero (8Byte)} = 14 \text{ byte}$

## Why is a bind system call required in socket programming?

- Otherwise if a port is not binded to any socket descriptor, we can use the same port and IP combination again to connect to a different destination.
- `bind()` assigns the address specified to by `addr` to the socket referred to by the file descriptor `sockfd`.
- `bind()` associates the socket with its local address [that's why server side binds, so that clients can use that address to connect to server.]

# sendto and recvfrom Functions

```
int sendto(int sockfd, char *buff, int length, int flags,  
struct sockaddr *to, socklen_t addrlen);
```

*to* argument is a socket address structure containing protocol address of where the data is to be sent

**sockfd** -- The socket descriptor.

**buff** --The pointer to the buffer containing the message to transmit.

**length** -- The length of the message in the buffer pointed to by the msg parameter.

**Flags**-- **MSG\_DONTROUTE** -- Don't use a gateway to send out the packet, only send to hosts on directly connected networks. This is usually used only by diagnostic or routing programs ;

**MSG\_DONTWAIT** - Enables nonblocking operation;

**MSG\_WAITALL** requests blocking until the entire number of bytes requested can be read.

**to** --The address of the target.

**addr\_len** --The size of the address pointed to by address.



# sendto and recvfrom Functions

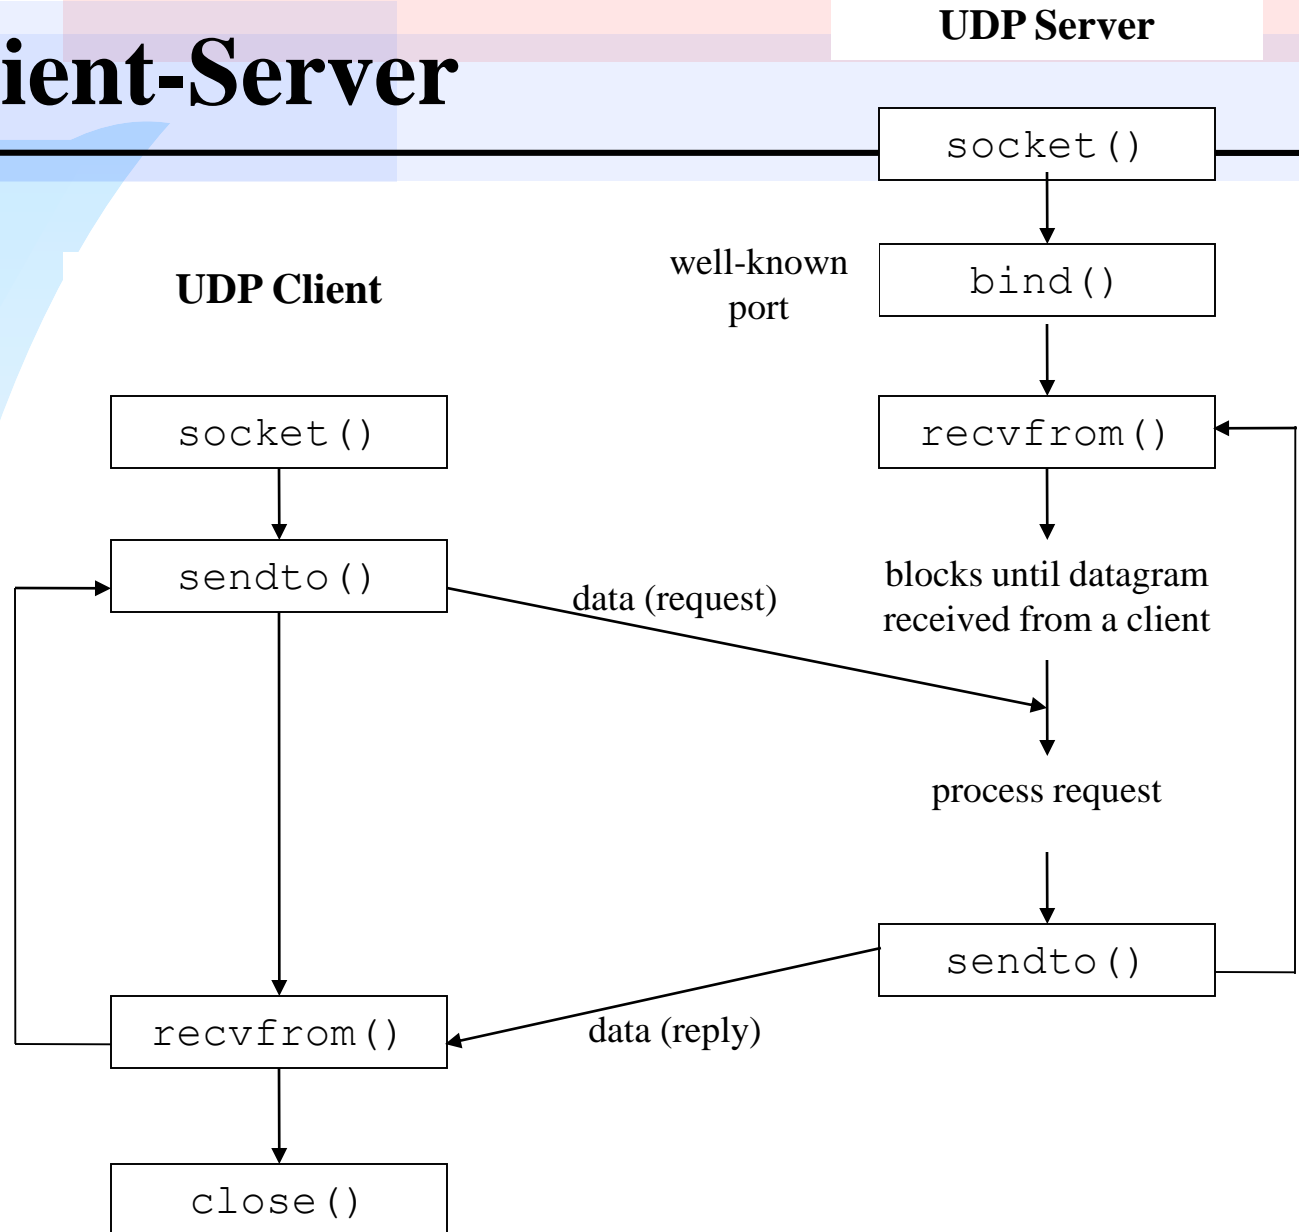
- A blocking system call is one that must wait until the action can be completed. `read()` would be a good example - if no input is ready, it'll sit there and wait until some is (provided you haven't set it to non-blocking, of course, in which case it wouldn't be a blocking system call)
- If there is not enough available buffer space to hold the socket data to be transmitted, and the socket is in blocking mode, `sendto()` blocks the caller until additional buffer space becomes available.

# sendto and recvfrom Functions

```
int recvfrom(int sockfd, char *buff, int nbytes, int  
flags, struct sockaddr *from, socklen_t *addrlen);
```

- Fills in the socket address structure pointed to by *from* with the protocol address of who sent the datagram
- If no incoming data is available at the socket, the recvfrom function blocks and waits for data to arrive according to the blocking rules

# UDP Client-Server



# Header Files

- `#include <sys/socket.h> // Core BSD socket functions and data structures.`
- `#include <netinet/in.h> // AF_INET and AF_INET6 address families and their corresponding protocol families PF_INET and PF_INET6.`
- `#include <arpa/inet.h> // Functions for manipulating numeric IP addresses.`
- `#include <netdb.h> // Name resolution`

# Example (UDP\_CLIENT)

- struct MSG  
{  
    int m\_txt[MAX\_TEXT];  
    int m\_seq;  
    int total\_no\_of\_packets;  
};
- struct MSG msg;
- struct sockaddr\_in sa;
- sock1 = socket(AF\_INET,SOCK\_DGRAM,0);
- bytes\_sent = sendto(sock1, (void\*)&msg, sizeof(struct MSG),0,(struct sockaddr\*) &sa, sizeof(struct sockaddr\_in) );


# Example (UDP\_SERVER)

- struct MSG  
{  
    int m\_txt[MAX\_TEXT];  
    int m\_seq;  
    int total\_no\_of\_packets;  
};
- struct MSG msg;
- int s\_length;
- struct sockaddr\_in sa,cli;
- sock1 = socket(AF\_INET,SOCK\_DGRAM,0);

# Example UDP\_SERVER

- `sa.sin_family = AF_INET;`
- `sa.sin_port = htons(PORT);`
- `sa.sin_addr.s_addr = htonl(INADDR_ANY);`
- `s_length=sizeof(sa);`
- `bind(sock1,(struct sockaddr *)&sa, s_length);`
- `rec=recvfrom(sock1,(void*)&msg,sizeof(struct MSG),0,(struct sockaddr *)&cli,sizeof(cli) );`

# THREE-WAY HANDSHAKING



**Passive Open:** The application on the server is passive. The application is listening, awaiting a connection

The application on the client makes a connection request to the server where the application is passive open.

The application on the client is said to be "**active open**".



# connect System Call

```
int connect(int sockfd, struct sockaddr *servaddr,  
            int addrlen);
```

- For connection oriented protocols the `connect` system call results in the actual establishment of a connection between the local and foreign system
- `sockfd` is a socket descriptor that was returned by the `socket` system call
- `servaddr` is a pointer to an address structure containing the address of the server
- `addrlen` is the length of the address structure

# connect()

```
#include <sys/socket.h>
int connect( int sockfd, const struct sockaddr * servaddr, socklen_t
addrlen );
```

*Returns: 0 if OK, -1 on error*

- e.g. connect()

```
    ■ struct sockaddr_in servaddr;
    ...
    servaddr.sin_family = AF_INET;
    servaddr.sin_addr.s_addr = htonl( "220.95.133.100" );
    servaddr.sin_port = htons( 7 );    // 7 for echo service

    if( connect( sockfd, (struct sockaddr *)&servaddr,
        sizeof(servaddr) ) < 0 )
    {
        ...
    }
    ...
```

# listen System Call

```
int listen(int sockfd, int backlog);
```

- This system call is used by a connection-oriented server to indicate that it is willing to receive connection
- Called after both the `socket` and `bind` system calls, and immediately before the `accept` system call
- The *backlog* argument specifies **how many connection requests can be queued by the system** while it waits for the server to execute the `accept` system call
- That number is only the **size of the connection queue**, where new connections wait for somebody to accept them. As soon as your application calls `accept()`, a waiting connection is removed from that queue. So, you can definitely handle more than 128 simultaneous connections because they usually only spend a short time in the queue

**Listen()**



**Connect()**

**Connect()**

**Connect()**

**Connect()**

**Queue**

**Accept()**

As soon as your application calls `accept()`, a waiting connection is removed from that queue.

recv SYN

send SYN+ACK



SYN Queue

recv ACK

Application



`accept()`

Read/write operation

**Buffer**



Accept Queue

# accept System Call

```
int accept(int sockfd, struct sockaddr *peer, int  
          *addrlen) ;
```

- Takes first connection request on the queue and creates another socket with the same property as *sockfd*
- Blocks the caller if no connection request is pending
- *peer* argument return the address of the connected peer client
- *addrlen* returns the length of the client address structure

# listen() and accept()

```
#include <sys/socket.h>
int listen( int sockfd, int backlog );
```

*Returns: 0 if OK, -1 on error*

```
int accept( int sockfd, struct sockaddr * cliaddr, socklen_t *
addrlen )
```

*Returns: non-negative descriptor if OK, -1 on error*

- e.g. listen()

- #define LISTENQ 5

- ...

- if( **listen**( **mysock**, **LISTENQ** ) < 0 ) { ... }

- while( 1 ){

- clilen = sizeof( cliaddr );

- if( ( clisockfd = **accept**( **mysock**, (struct sockaddr \*)&**cliaddr**,  
                          &**clilen** ) ) < 0 ) { ... }

- ...

- }

# recv() and send()

```
#include <sys/types.h>
#include <sys/socket.h>
int recv( int sockfd, void * msg, size_t len, int flags );
int send( int sockfd, const void * msg, size_t len, int flags );
```

*Both return: number of bytes read or written if OK, -1 on error*

**msg**: message you want to send

**len**: length of the message

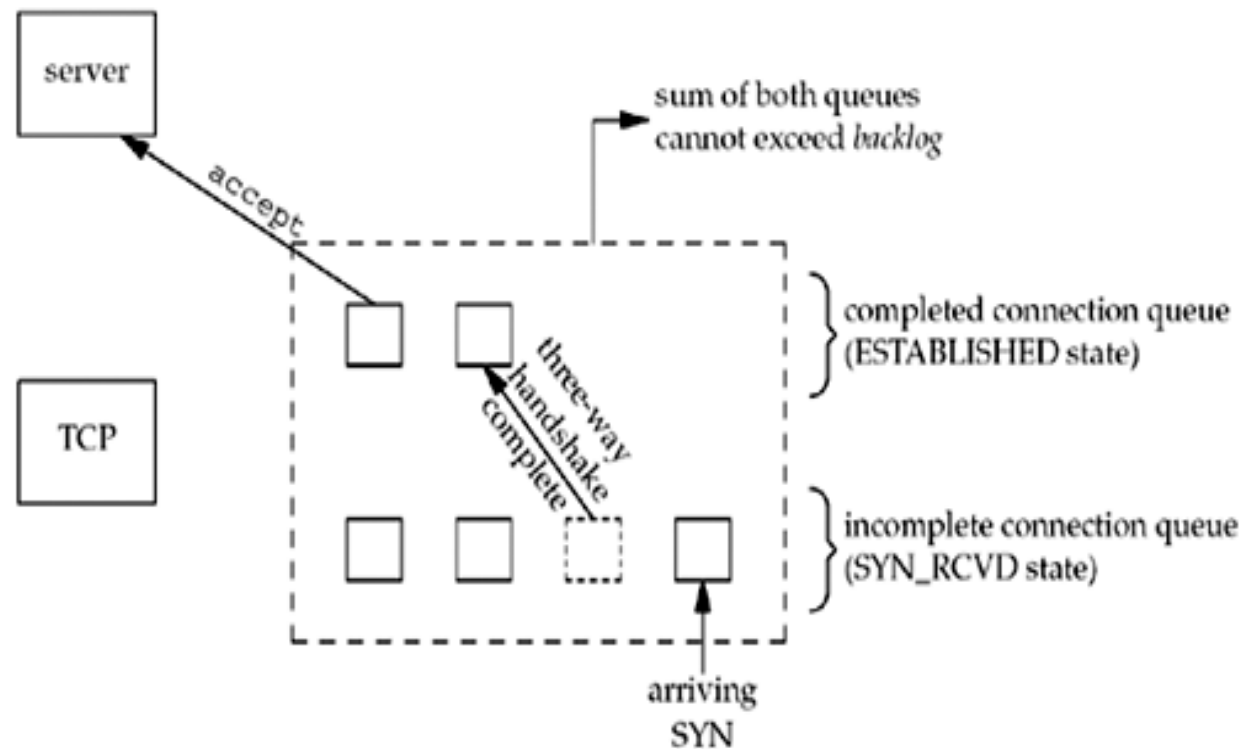
**flags** := 0

**returned**: the number of bytes actually sent or received

- e.g. send() and recv()


```
■ if( ( n = send( sockfd, buff, len, 0 ) ) != len ) { ... }
...
if( ( n = recv( sockfd, buff, BUFMAX, 0 ) ) < 0 ) { ... }
```

The two queues maintained by TCP for a listening socket.

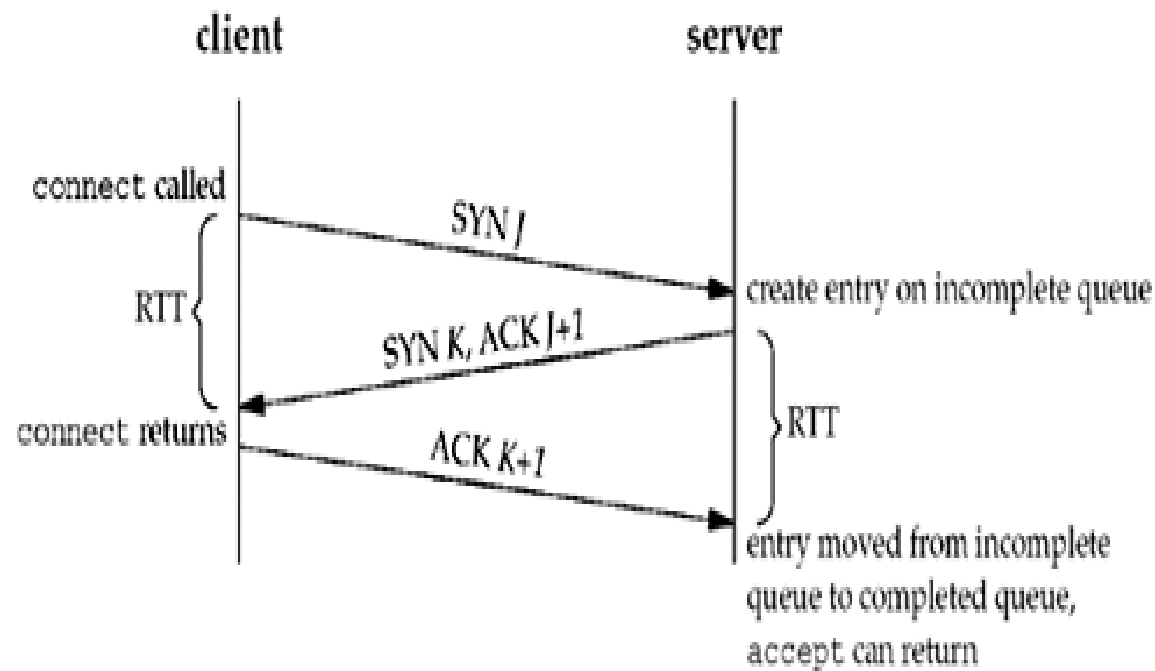




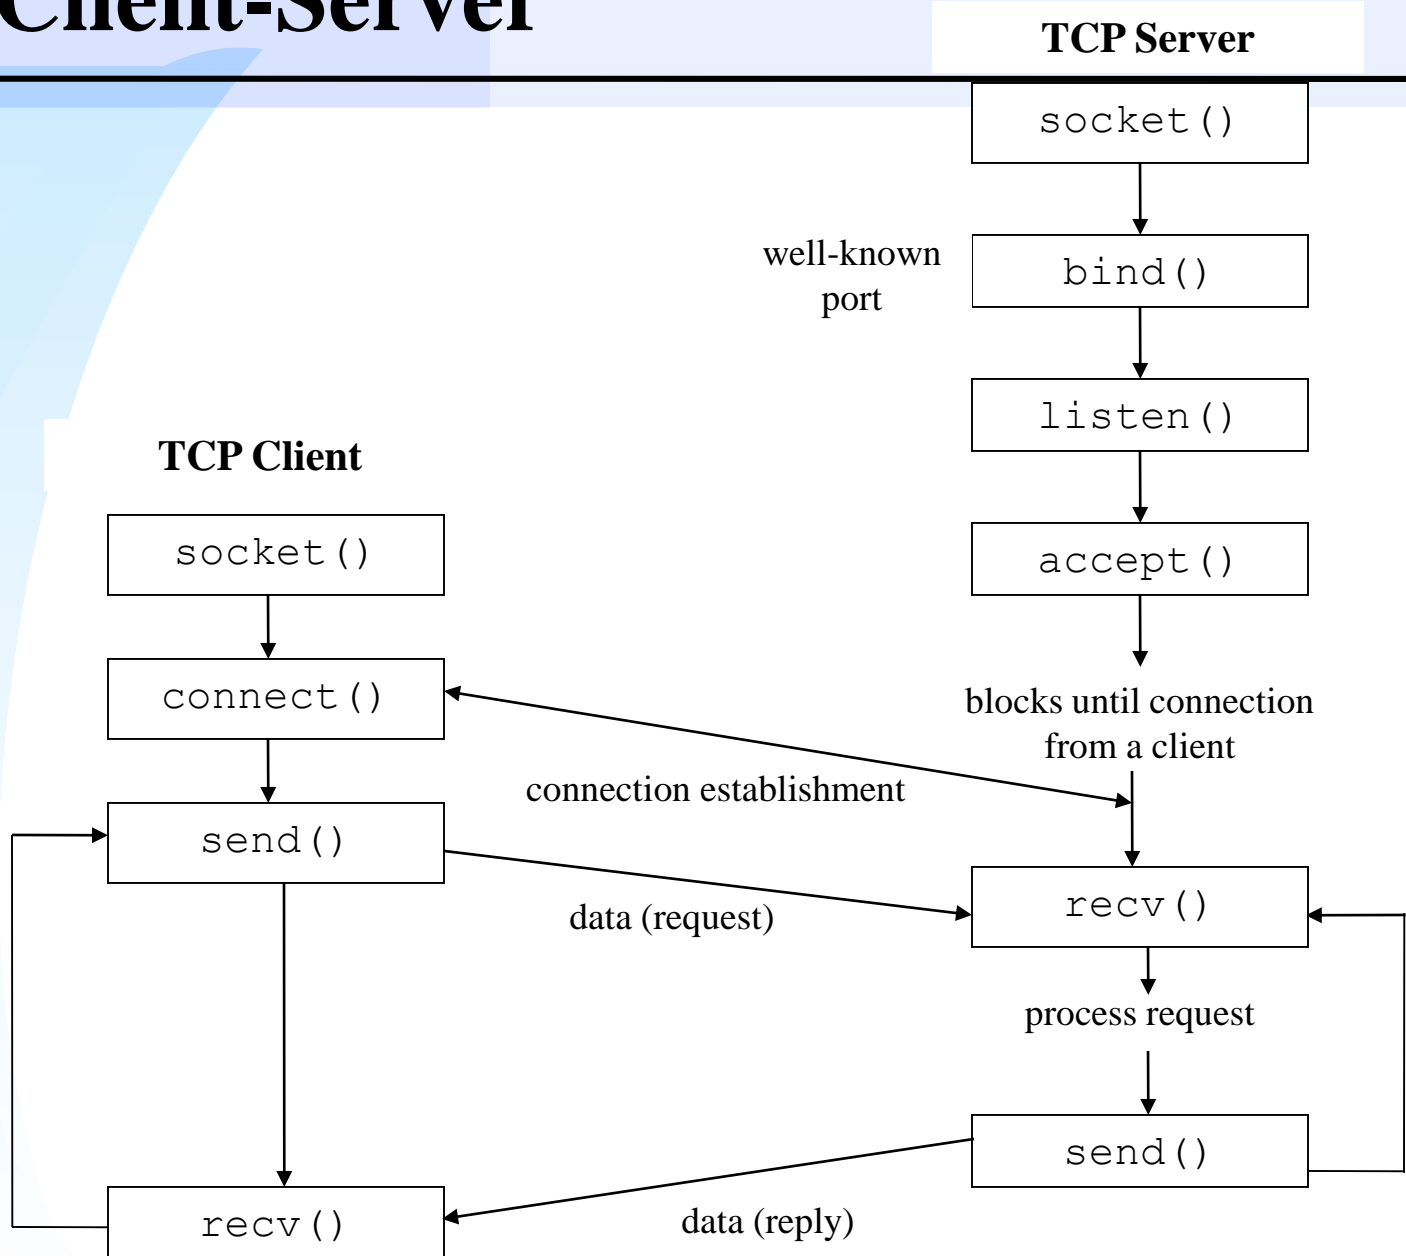
# THREE-WAY HANDSHAKING



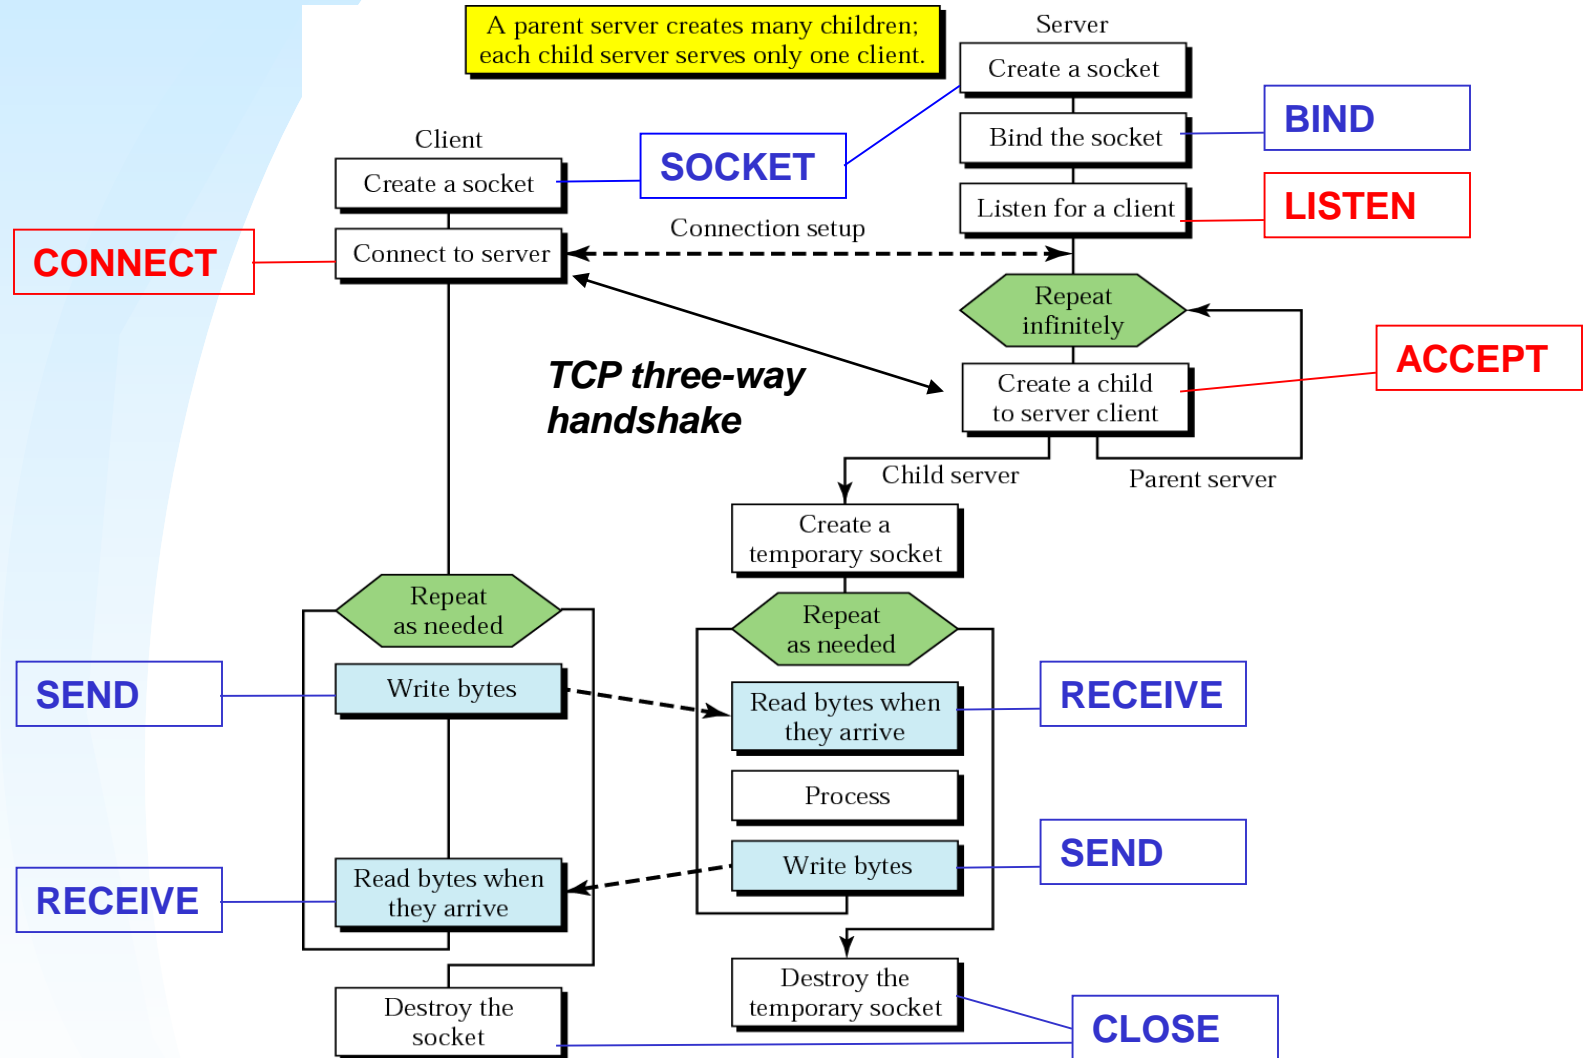
## TCP three-way handshake and the two queues for a listening socket.



# TCP Client-Server



# Client+server: connection-oriented



Concurrent server

# TCP SERVER

- struct sockaddr\_in server\_addr, client\_addr;
- int sin\_size;
- if ((sock = socket(AF\_INET, SOCK\_STREAM, 0)) == -1)  
    perror("Socket");  
    exit(1);  
}
- server\_addr.sin\_family = AF\_INET;
- server\_addr.sin\_port = htons(5000);
- server\_addr.sin\_addr.s\_addr = INADDR\_ANY;
- bzero(&(server\_addr.sin\_zero), 8);
- if (bind(sock, (struct sockaddr \*)&server\_addr, sizeof(struct sockaddr)) == -1)  
{  
    perror("Unable to bind");  
    exit(1);  
}

# TCP SERVER

- `if (listen(sock, 5) == -1)`  
    {  
        `perror("Listen");`  
        `exit(1);`  
    }
- `sin_size = sizeof(struct sockaddr_in);`
- `connected = accept(sock, (struct sockaddr *)`  
    `&client_addr, &sin_size);`
- `char send_data[1024],recv_data[1024];`
- `gets(send_data);`
- `send(connected, send_data,strlen(send_data), 0);`

# TCP CLIENT

- struct sockaddr\_in server\_addr;
- if ((sock = socket(AF\_INET, SOCK\_STREAM, 0)) == -1)  
{  
    perror("Socket");  
    exit(1);  
}
- server\_addr.sin\_family = AF\_INET;
- server\_addr.sin\_port = htons(5000);
- server\_addr.sin\_addr = inet\_addr("server IP Address"); //inet\_addr("192.168.148.12")
- bzero(&(server\_addr.sin\_zero),8);
- if (connect(sock, (struct sockaddr \*)&server\_addr, sizeof(struct sockaddr)) == -1)  
{  
    perror("Connect");  
    exit(1);  
}

# TCP CLIENT

---

- `char recv_data[1024];`
- `bytes_recieved=recv(sock,recv_data,1024,0);`



# References

---

- Unix Network Programming, Volume 1, Second Edition
  - W Richard Stevens
  - Pearson Education



